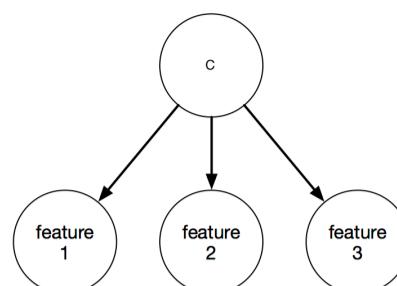


Project 2: Naïve Bayes



"All models are wrong, but some are useful."
-- George E.P. Box

Introduction

Aðalbrandr is visiting St. Louis from Norway and has been having the hardest time distinguishing boys and girls because of the weird American names like Jack and Jane. This has been causing lots of problems for Aðalbrandr when he goes on dates. When he heard that Washington University has a Machine Learning class, he asked that we help him identify the gender of a person based on their name to the best of our ability. In this project, you will implement Naïve Bayes to predict if a name is male or female. But before you get started, you need to create a GitHub Classroom team and clone the project2 repository.

The code for this project (project2) consists of several files, some of which you will need to read and understand in order to complete the assignment, and some of which you can ignore.

Files you'll edit:

naivebayesPXY.py	Computes $P(X Y)$
naivebayesPY.py	Computes $P(Y)$
naivebayes.py	Computes log probability $\log P(Y X = x_1)$
naivebayesCL.py	Turns the Naïve Bayes output into a linear classifier
classifyLinear.py	Makes predictions with a linear classifier.

Files you might want to look at:

name2features.py	Converts a name to a feature vector. You do not need to edit this, but you may want to!
genTrainFeatures.py	Calls name2features.py.
whoareyou.py	Allows you to test your linear classifier interactively.
example_tests.py	Describes several unit tests to find obvious bugs in your implementation.
boys.train	The training examples for boys names.
girls.train	The training examples for girls names.

How to submit: You can commit your code through the command line with git and submit on Gradescope either in a zip file or through Github. If the project is submitted before the initial deadline passes, you will receive information and a score for the performance evaluation (only once the deadline is reached). However, the autograder will not reveal any information on how your code performed for any projects submitted during the three day extension period. You can submit your project as many times as you want but the final submission score will count for your grade. If you submitted by the initial deadline and would like to improve your performance score, you can submit again during the extension period.

Grading: Your code will be autograded for technical correctness. Please *do not* change the names of any provided functions or classes within the code, or you will wreak havoc on the autograder. However, the correctness of your implementation -- not the autograder's output -- will be the final judge of your score. If necessary, we will review and grade assignments individually to ensure that you receive due credit for your work.

PYTHON Version in Autograder: The autograder uses PYTHON 3.6. To rule out any incompatibilities of different versions we recommend to use any version of PYTHON 3.6 or newer for the implementation projects.

Regrade Requests: Use Gradescope for regrade requests.

Academic Dishonesty: We will be checking your code against other submissions in the class for logical redundancy. If you copy someone else's code and submit it with minor changes, we will know. These cheat detectors are quite hard to fool, so please don't try. We trust you all to submit your own work only; *please* don't let us down. If you do, we will pursue the strongest consequences available to us.

Getting Help: You are not alone! If you find yourself stuck on something, contact the course TAs for help. Office hours and [Piazza](#) are there for your support; please use them. If you can't make our office hours, let us know and we will schedule more. We want these projects to be rewarding and instructional, not frustrating and demoralizing. But, we don't know when or how to help unless you ask.

Of boys and girls

Take a look at the files [girls.train](#) and [boys.train](#). For example with the unix command

```
cat girls.train
```



...
Addisyn
Danika

Emilee
Aurora
Julianne
Sophia
Kaylyn
Litzy
Hadassah

Believe it or not, these are all more or less common girl names. The problem with the current file is that the names are in plain text, which makes it hard for a machine learning algorithm to do anything useful with them. We therefore need to transform them into some vector format, where each name becomes a vector that represents a point in some high dimensional input space.

name2features.py reads a name, for example one in girls.train (or in boys.train) and converts it into a 128-dimensional feature vector. Can you figure out what the features are? Understanding how these features are constructed and represented will help you with your naivebayesPY.py implementation and also later on in the competition.

We have provided you with a Python function genTrainFeatures.py, which calls this Python script, transforms the names into features and loads them into Python. You can call

```
[x,y]=genTrainFeatures()
```

to load in the features and the labels of all boys and girls names.

The Naïve Bayes Classifier

The Naïve Bayes classifier is a linear classifier based on Bayes Rule. The following questions will ask you to finish these functions in a pre-defined order.

Remember: as a general rule, you should avoid tight loops at all cost.

- (a) Estimate the class probability $P(Y)$ in naivebayesPY.py . This should return the probability that a sample in the training set is positive or negative, independent of its features.
- (b) Estimate the conditional probabilities $P(X|Y)$ in naivebayesPY.py . Use a **categorical distribution** as model and make sure you understand why this is the appropriate model for our features generated in name2features.py. Make sure you understand the feature representation generated by name2features.py. Use array operations in numpy rather than for loops! This will return the necessary components of the probability vectors/CPTs for all features given each class label.
- (c) Solve for the log ratio, $\log\left(\frac{P(Y=1|X)}{P(Y=-1|X)}\right)$, using Bayes Rule. Implement this in naivebayes.py.
- (d) Naïve Bayes can also be written as a linear classifier. Implement this in naivebayesCL.py
- (e) Implement classifyLinear.py that applies a linear weight vector and bias to a set of input vectors and outputs their predictions. (You can use your answer from the previous project.) You can now test your training error with

```
>> [x,y]=genTrainFeatures();
>> [w,b]=naivebayesCL(x,y);
>> predss=classifyLinear(x,w,b);
>> trainingerror=np.sum(preds~y)/(y.shape[1])
trainingerror = 0.XX...
```

Once your Naïve Bayes code is implemented, you can try the function whoareyou.py to test names interactively.

In a script:

```
xTr,yTr = genTrainFeatures()
w, b = naivebayesCL(xTr, yTr)
whoareyou(w, b)
```

In the terminal:

```
Who are you>David
David, I am sure you are a nice boy.
Who are you>Anne
Anne, I am sure you are a nice girl.
Who are you>byebye
```

Hints

Tests. To test your code you can run example_tests.py, which describes and partially implements several unit tests. Feel free to implement your own tests to better debug and test your implementation. Those tests are a subset of what we will use in the autograder to grade your submission.

70% of the grade for your project2 submission will be assigned based on the correctness of your implementation.

Feature Extraction (Quality Evaluation)

30% of the grade for your project2 submission will be assigned by how well your classifier performs on a secret test set. If you want to improve your classifier modify name2features.py. The autograder will use your Python script to extract features and train your classifier on the names training set.